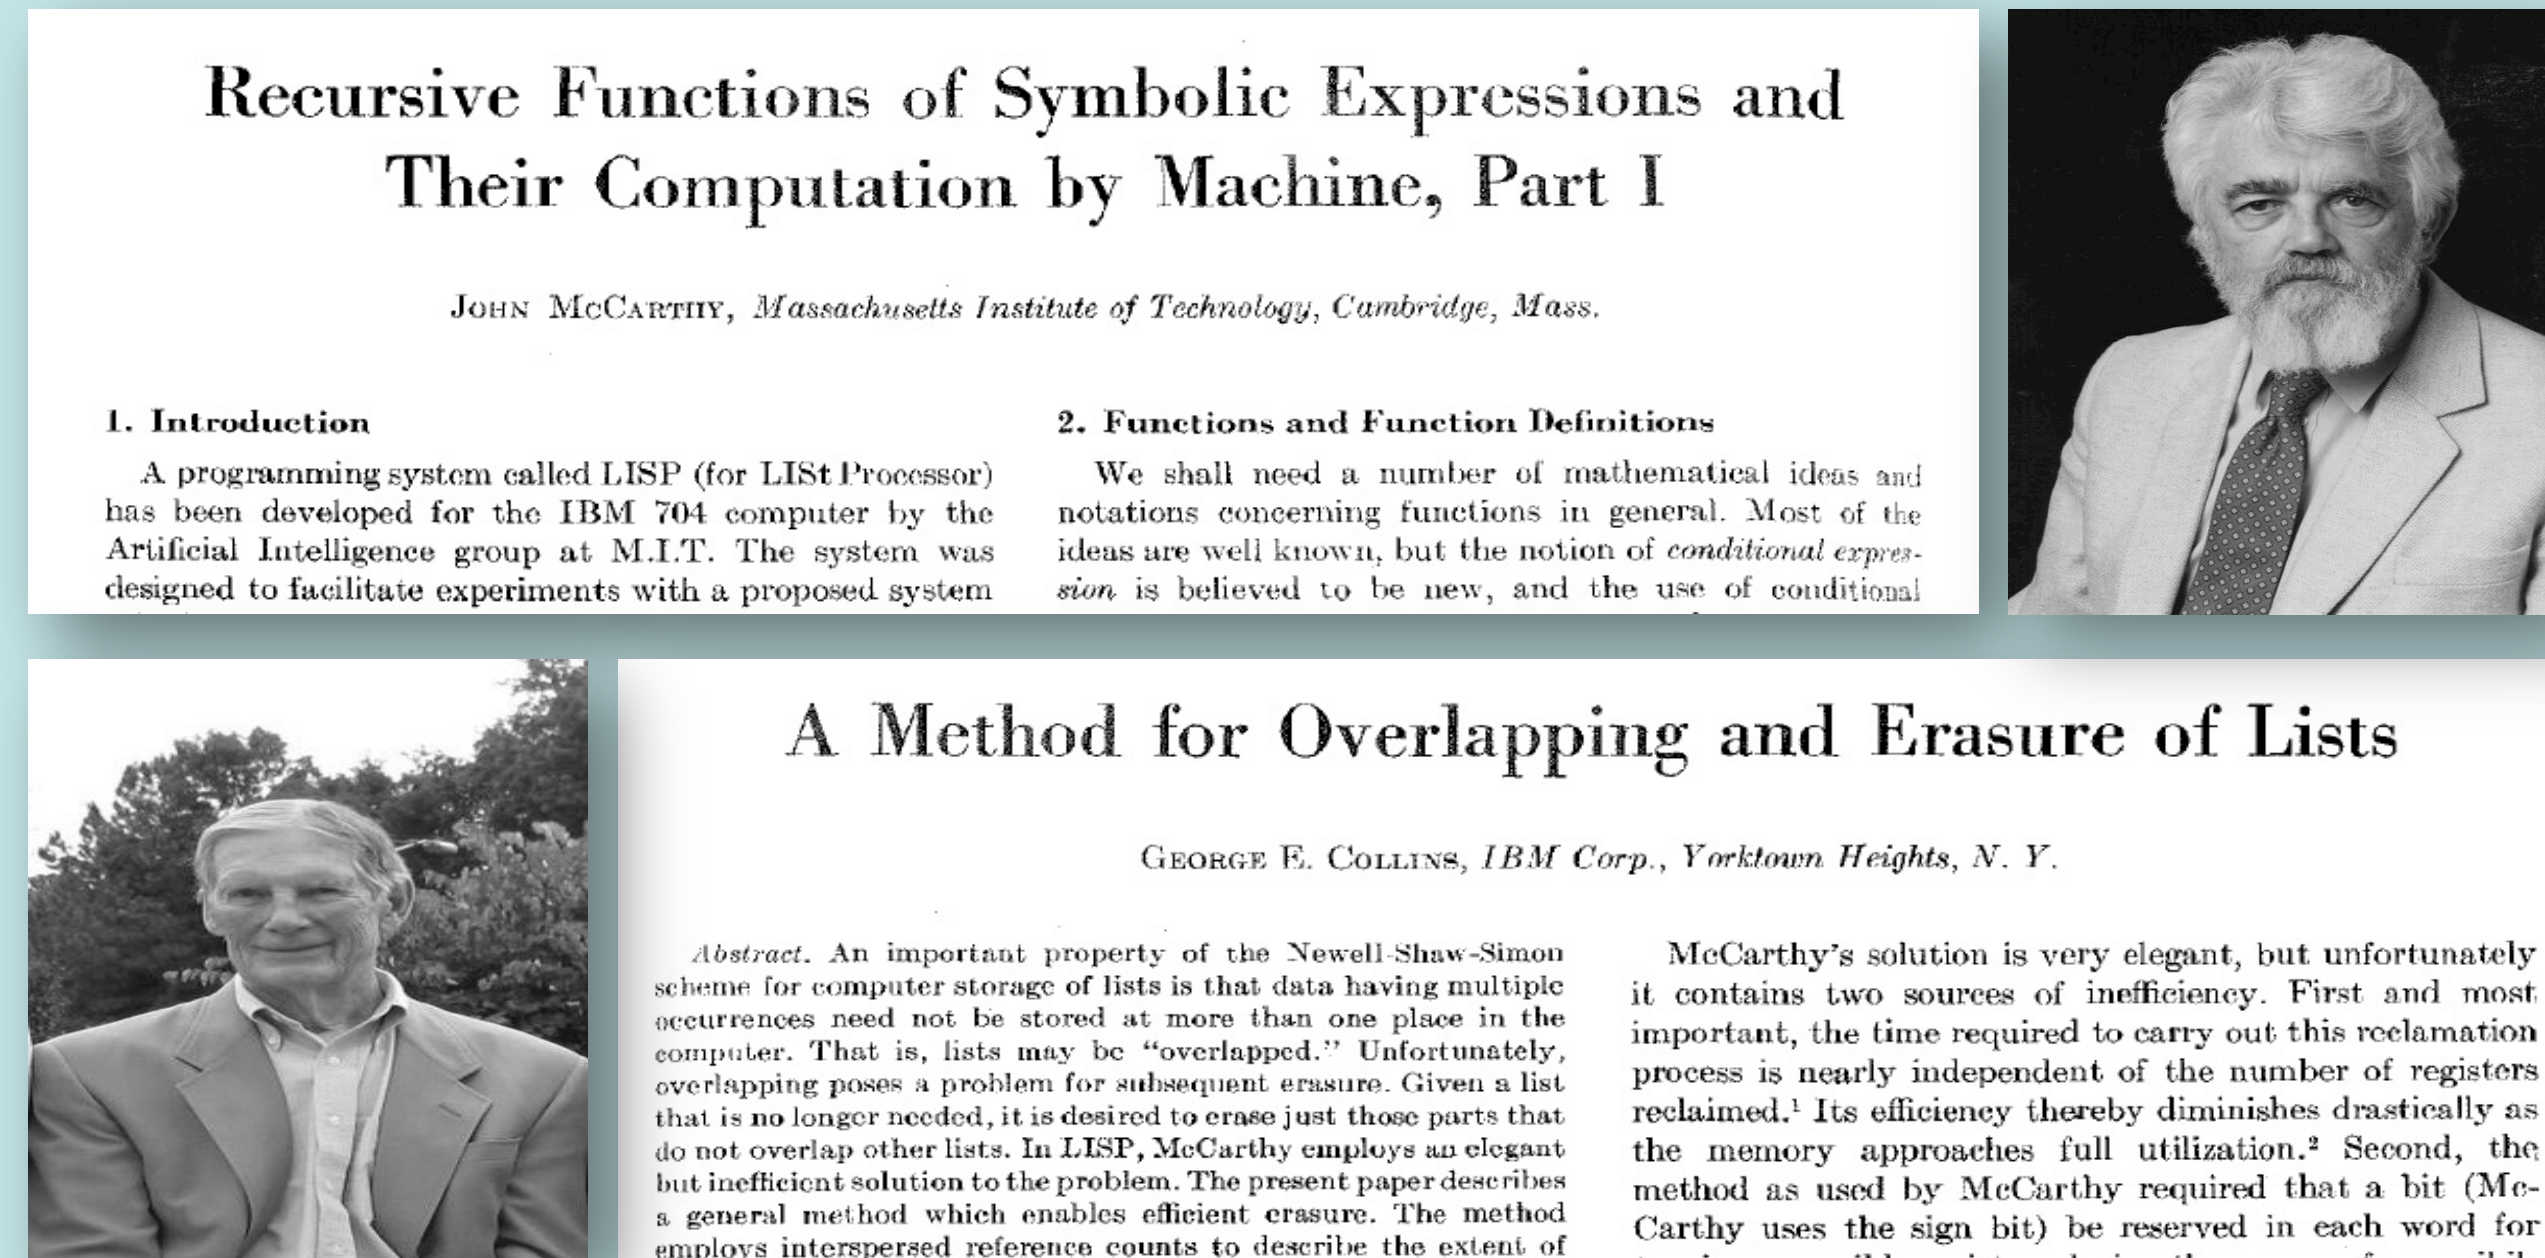


Down for the Count? Getting Reference Counting Back in the Ring

By Rifat Shahriyar[†], rifat.shahriyar@anu.edu.au. Supervisors: Steve Blackburn[†] and Daniel Frampton^δ

Garbage collection (GC) is Ubiquitous

- Born 52 years ago



- Two ideas underpin large literature:
 - **Tracing** [McCarthy60]
 - **Reference Counting** [Collins60]
- However
 - ✓ Tracing used in all high performance GCs
 - ✗ Reference counting (RC) only in non-performance critical settings

Reference Counting vs. Tracing

- **Advantages**
 - ✓ Immediate
 - ✓ Incremental (Reclaim as-you-go)
 - ✓ Object-local
 - ✓ Overhead distributed
 - ✓ Very simple (Trivial implementation for naïve RC)
- **Disadvantages**
 - ✗ Maintain count (Time and space overheads)
 - ✗ Cycles (Can't be collected)
 - ✗ Complex (High performance implementation about as complex as tracing)

The Challenge


- ✓ One of the two fundamental GC algorithms
- ✓ Many advantages
- ✗ Neglected by all performance-conscious VMs
- ✗ High performance RC is 30% slower than mark-sweep (MS)

Can we get RC back in the ring?

[†] Computer Systems Group, Research School of Computer Science, College of Engineering & Computer Science, Australian National University, ^δ Microsoft

Storing the Count


- **Space**
 - Dedicated word (32 bits) per object
 - Steal bits from each object's header
- **Findings**
 - Max count < 8 in most cases
 - Very few overflows (The percentage of stuck objects is very low)



- **Design: Handling stuck objects**
 - Auxiliary data structure to store count of the overflowed objects
 - Ignore them let backup tracing collect stuck objects
 - Restore them let backup tracing restore stuck counts

Maintaining the Count

- **Types**
 - Deferred RC ignores changes to stacks and register
 - Coalescing RC ignores many changes to heap
- **Findings**
 - New objects are the source of most incs and decs
 - Survival ratio is low
 - New objects a fruitful focus for optimization



Handling of New Objects

- **Existing**
 - Implicitly dirty (marked as dirty and inc enqueued for next collection)
 - Implicitly live (initial count of one and dec enqueued for next collection)
- **New**
 - Implicitly clean (lazily dirty at collection time, non-surviving never processed)
 - Implicitly dead (lazily increment at collection time, available to free list if no inc)

Performance Improvement

- ✓ Our improved RC is 24% faster than the standard RC
- ✓ Standard RC was 30% slower than MS, but our improved RC performs the same
- ✓ Improved RC is only 2% slower than URC and 3% slower than Immix
- ✓ It performs same as Sticky MS but 10% slower than Sticky Immix (current project)

RC is back in the ring

