

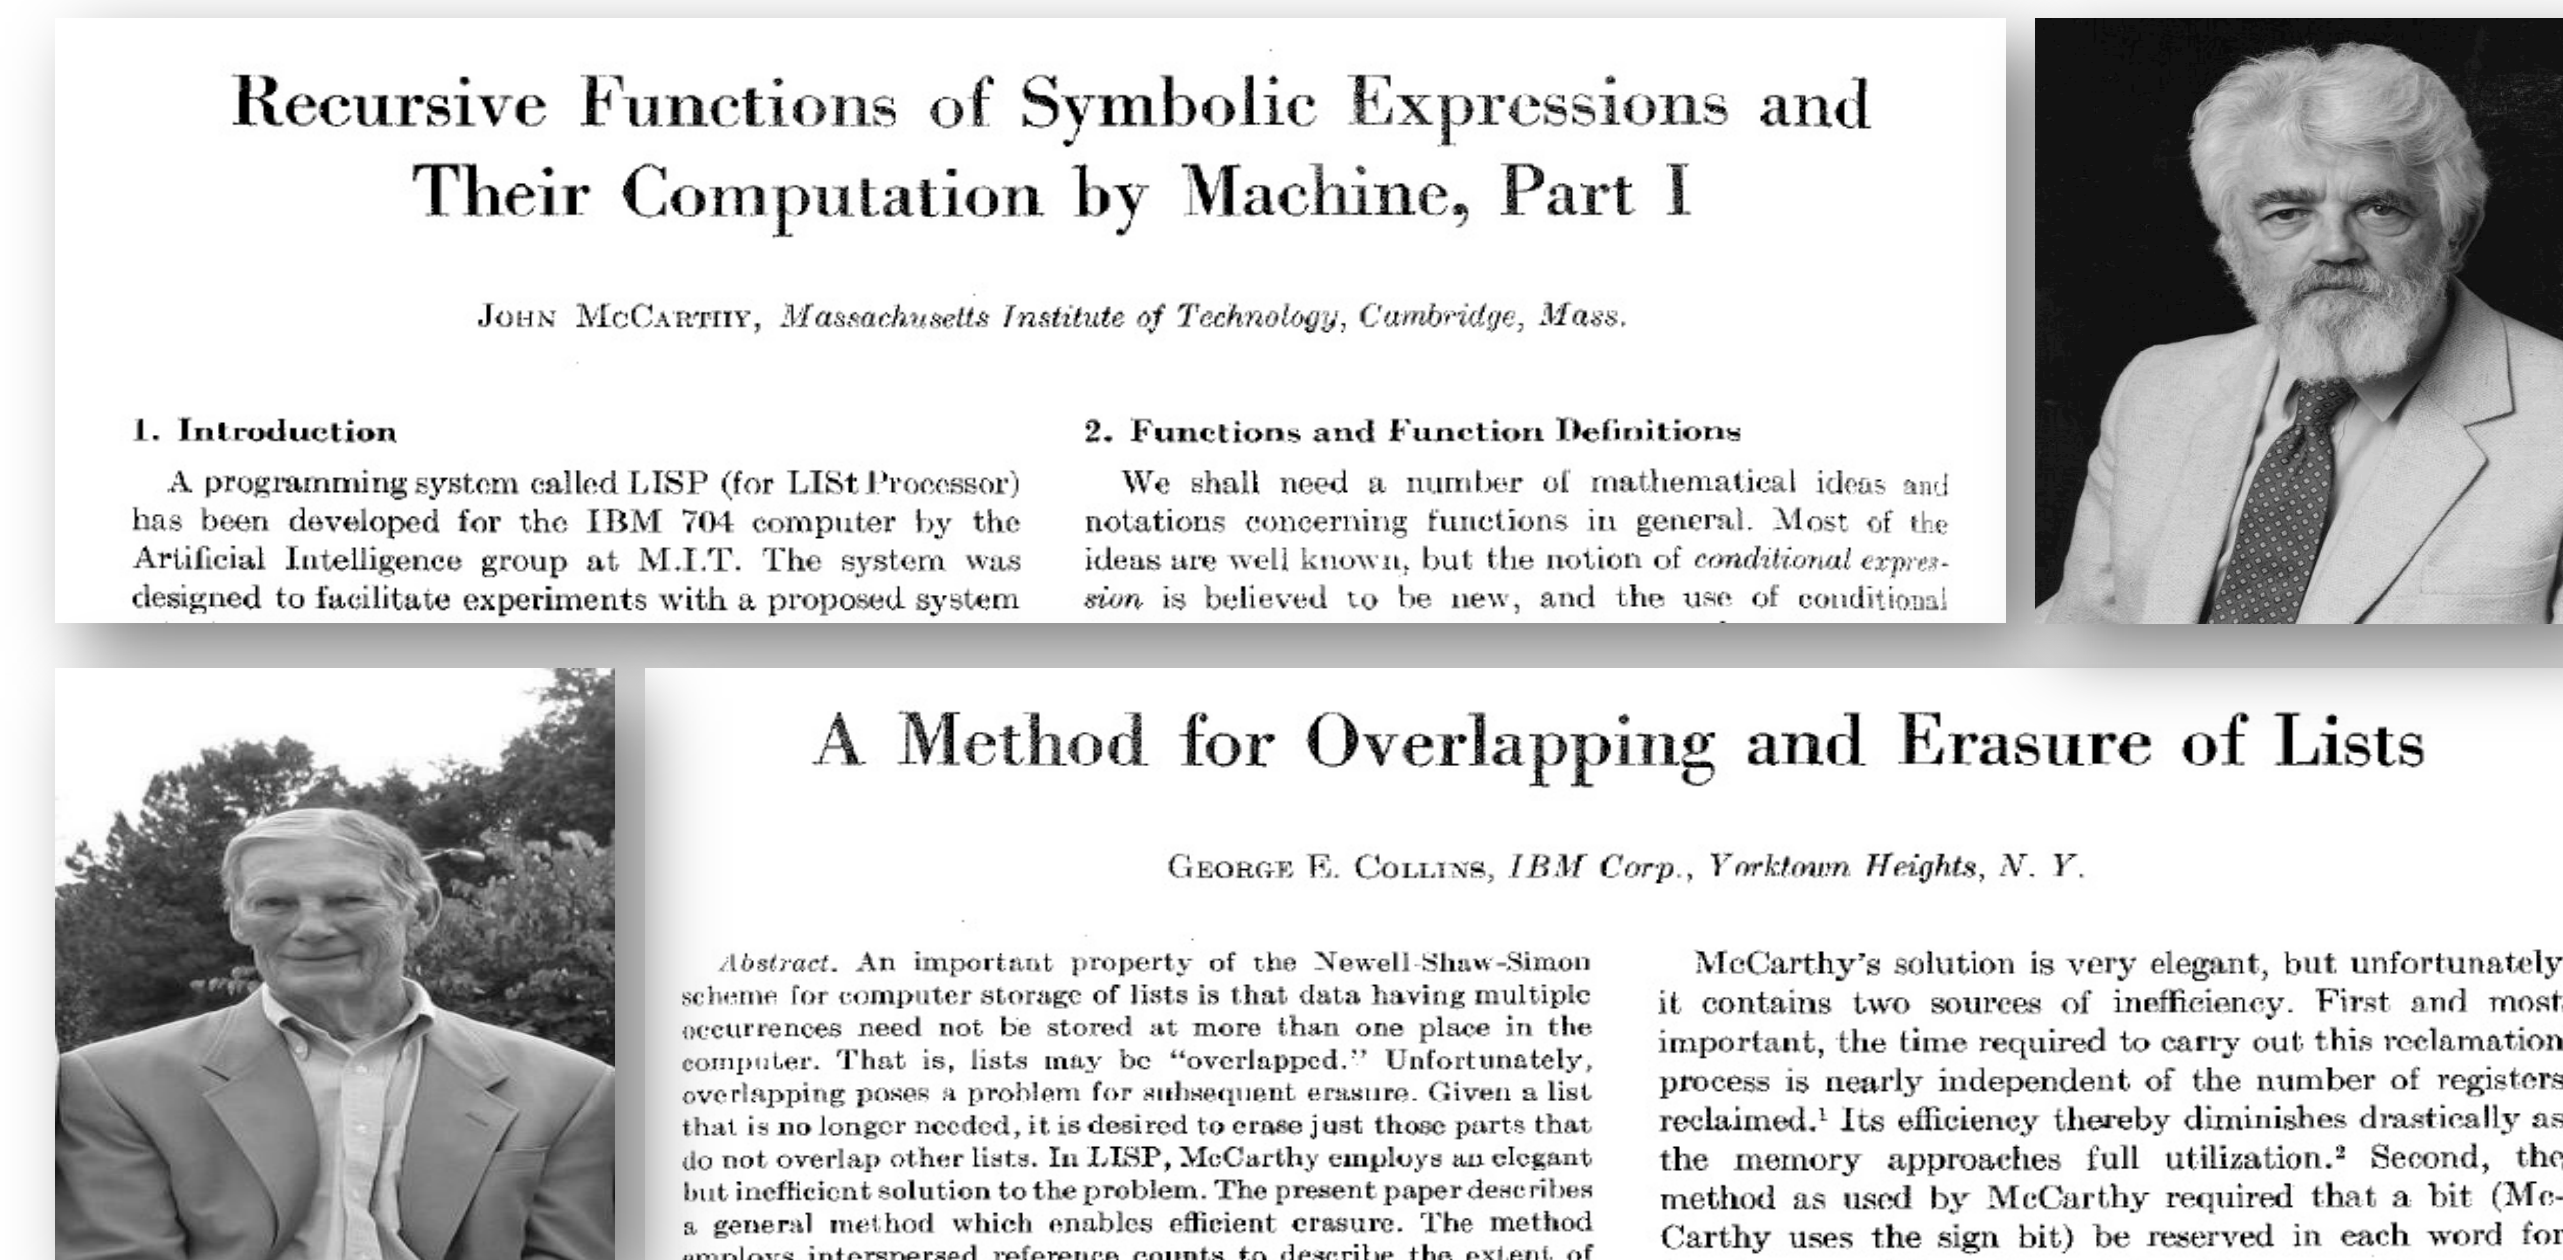
# Taking Off the Gloves with Reference Counting Immix



Rifat Shahriyar<sup>†</sup>, Steve Blackburn<sup>†</sup>, Xi Yang<sup>†</sup> and Kathryn McKinley<sup>δ</sup>

## Garbage collection (GC) is Ubiquitous

- Born 53 years ago



- Two ideas underpin large literature:
  - Tracing [McCarthy60]
  - Reference Counting [Collins60]
- However
  - ✓ Tracing used in all high performance GCs
  - ✓ Reference counting (RC) has interesting advantages
  - ✗ Reference counting only in non-performance critical settings

## Status of Reference Counting

- **High performance reference counting**
  - ✓ Significantly faster than naïve RC
  - ✗ 30% slower than MS (well tuned simple tracing)
  - ✗ 40% slower than GenImmix (production collector in Jikes RVM)
- **Reference counting was improved [ISMM12]**
  - ✓ Deferred and coalesced limited bit RC with new object optimization
  - ✓ Performs same as MS
  - ✗ But 10% slower than GenImmix

## Allocator

- **Contiguous allocator**
  - ✓ Better cache locality
  - ✓ Fewer instructions per allocation
- **Free list allocator**
  - ✓ Suitable for RC
  - ✗ Poor cache locality
  - ✗ Higher instructions per allocation
  - ✗ Suffers from both internal and external fragmentation

## Motivating Analysis

GC	Allocator	Mutator time	Instruction retired	Cache miss	Mutator locality
Immix	Contiguous	1.00	1.00	1.00	✓
MS	Free list	1.09	1.07	1.27	✗
RC	Free list	1.12	1.12	1.31	✗
SS	Contiguous	1.01	1.00	0.97	✓

## Contributions

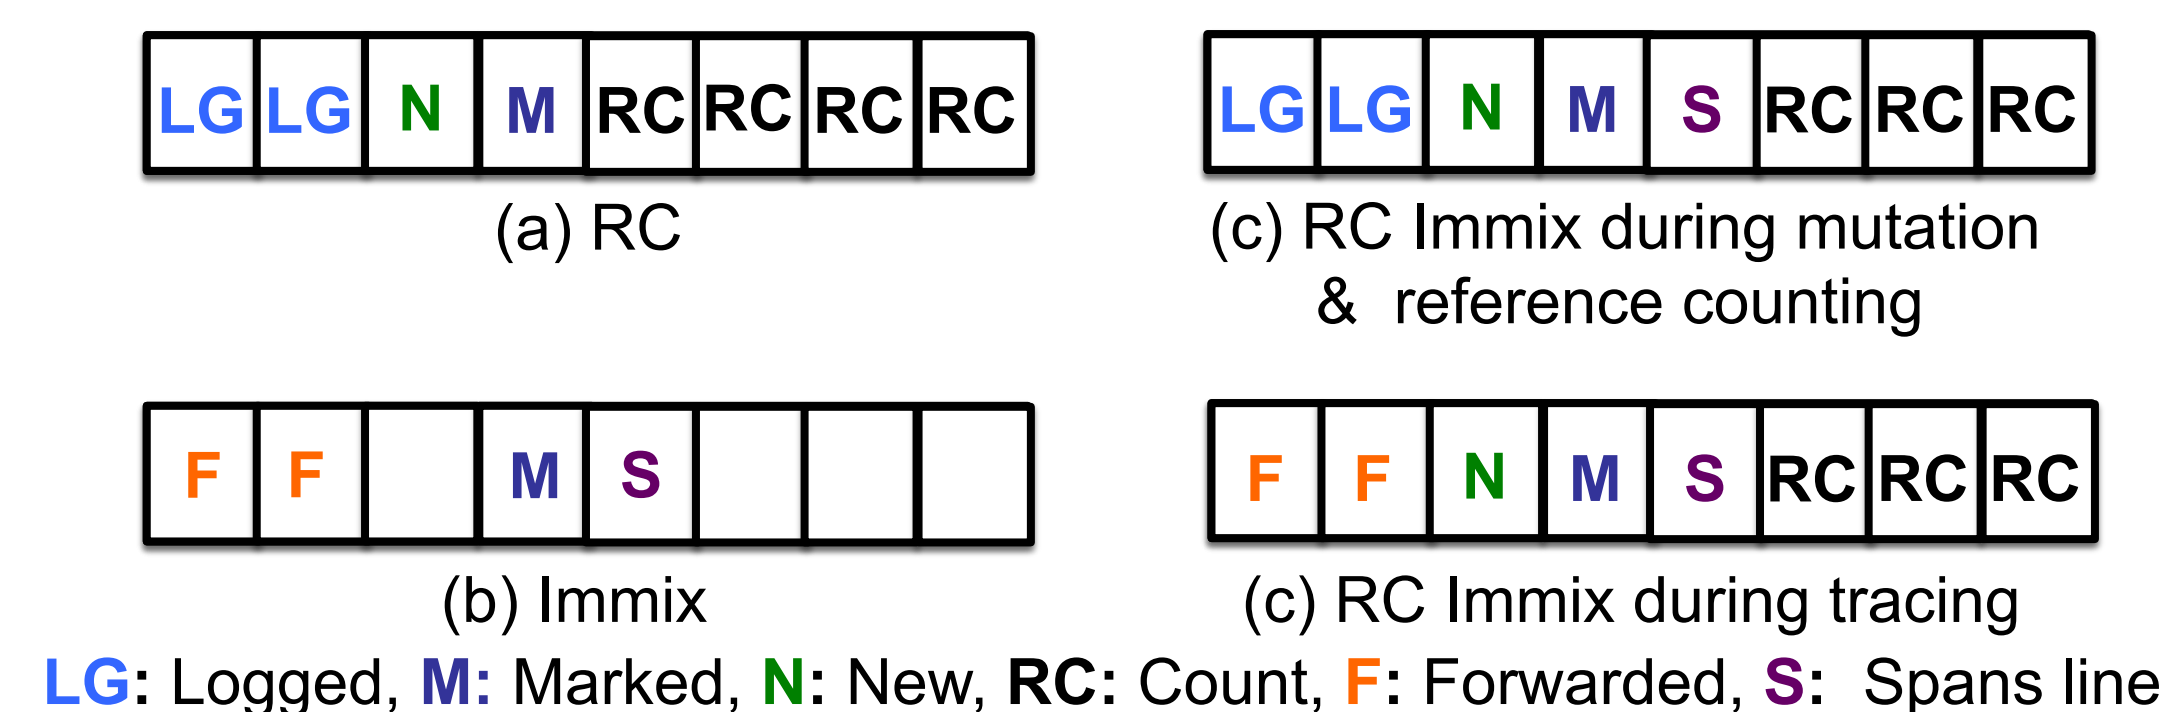
- ✓ Identify heap organization as performance bottleneck for RC
- ✓ Merge RC with Immix - RCImmix
- ✓ Eliminate fragmentation by integrating copying with RC
- ✓ RCImmix achieved great performance, 3% faster than fastest production

## Challenges of RCImmix

- Adapt Immix line/block reclamation strategy to RC context
- Share limited header bits to satisfy both RC and Immix
- Defragment in RC context to eliminate fragmentation






## How RCImmix works

- **Immix heap organization**
  - Contiguous allocation into regions (lines and blocks)
  - Mark objects and their region, unmarked regions can be freed
- **Reference counting collection**
  - Reference count for each object, live object count for each line
  - Collect lines with no live objects
- **Cycle collection**
  - Mark objects and their lines, sweep to collect unmarked lines
  - Restore stuck object counts and correct incorrect line counts
  - Sweep dead lines instead of sweep dead objects
- **Defragmentation**
  - Proactively copies surviving new objects with bounded copy reserve
  - Copy reserve using line survival rate without any overhead
  - Reactively with cycle collection based on some statistics and threshold
  - Both copies opportunistically and stops when available space exhausted
- **Header Bits**



## Performance Improvement

- ✓ RCImmix is 12% faster than RC at moderate (2x) heap size
- ✓ RCImmix outperforms the fastest production (GenImmix) by 3% at 2x
- ✓ RCImmix matches GenImmix at 1.3x and outperforms from 1.4x



## Future Opportunities

- Root Coalescing – unnecessary increment and decrement for unchanged roots
- Conservative Stack Scanning – enable to use RCImmix instead of naïve RC

## Summary

- ✓ RCImmix, a new GC by combining RC and Immix, outperforms fastest production
- ✓ Transforms RC into a serious alternative to meet high performance objectives for GC languages

<sup>†</sup> Australian National University <sup>δ</sup> Microsoft Research



Australian  
National  
University

Microsoft®  
Research